

Benjamin Johnston

(614) 843-3974 | bejaunst@gmail.com | linkedin.com/in/bejaunst | github.com/cetio | cetio.github.io

TECHNICAL SKILLS

Languages: C#, C, D, Python, JavaScript, TypeScript, Ruby, SQL

AI & LLM: RAG, LLM Evaluation, Prompt Engineering, AI Agents, Tool Calling, Structured Output, Embeddings, Vector Search, MCP, OpenAI API, Claude API

Infrastructure: Docker, CI/CD, GitHub Actions, REST APIs, OAuth 2.0, JSON Schema, SDK Design

EXPERIENCE

Software Engineer

Jul 2021 – Present

Independent client engagements

- Authored Intuit, the first native LLM SDK for D, mapping chat completions, structured output, and tool calls onto typed structures with configurable-precision embeddings.
- Developed Warbler, an AI-assisted grading platform contracted for high-school use, with LLM subagent architecture, JSON schema enforcement, and a D CLI backend with retry-with-validation loops (~95% grading accuracy, halved instructor workload).

AI Operations Engineer

Nov 2025 – Present

Handshake

- Evaluated 4,000+ multimodal LLM responses against structured rubrics for correctness and reasoning quality (LLM-as-a-judge evaluation), earning the Handshake Model Validation 2 — Expert credential.
- Eliminated 95% of manual annotation effort across 400+ runtime hours by building AutoHandshake!, a Ruby/JavaScript Selenium RPA suite with unattended execution, UI-stall recovery, and multi-tab orchestration.
- Mentored contractors on rubric interpretation and evaluation workflows, improving consistency and reducing rework from misaligned quality standards.

AI Evaluation Engineer

Dec 2025 – Jun 2026

Snorkel AI

- Engineered adversarial red-team tasks for Claude and ChatGPT across 2 contracted projects, applying strict schema constraints to surface model blindspots and reasoning failure modes.
- Validated training-set integrity via Docker-based evaluation sandboxes, catching defects pre-release and feeding red-team findings into model alignment.

SELECTED PROJECTS

Intuit | *LLM SDK, D* | github.com/cetio/intuit

- Shipped the first native LLM SDK for D with 3 provider integrations (OpenAI, Claude, Qwen), OpenAI-compatible routing for OpenRouter and LiteLLM, and 23 example applications from RAG to multi-agent debates.
- Exposed native D functions as model tools via automatic JSON schema generation for type-safe tool use, with configurable-precision embeddings.

Akashi + Chemica | *RAG Pipeline, D* | github.com/cetio/chemica

- Constructed a RAG pipeline aggregating compound data from 6 sources (PubChem, Wikipedia, PsychonautWiki, Erowid, PubMed, PMC) with 3D molecular visualization, structural similarity matching, and local embeddings.
- Extended the pipeline with LLM-powered source synthesis across 5 providers, enabling cross-source conflict resolution via local embeddings.

Insurgent | *D, JavaScript* | github.com/insurgent-con

- Delivered a contracted real-time companion system (API + client) for a content creator's community, working alongside Twin Harbor Interactive developers.
- Built a full-stack system rendering 10,000+ tiles and 5,000+ units via WebSocket feeds in under a second, with Selenium-driven live data capture and cached match state.

EDUCATION & CERTIFICATIONS

Columbus State Community College

Expected Dec 2026

AAS, Software Development

Handshake Model Validation 2 — Expert | MS Career Essentials Generative AI (NASBA)