

Benjamin Johnston

(614) 843-3974 | bejaunst@gmail.com | linkedin.com/in/bejaunst | github.com/cetio | cetio.github.io

TECHNICAL SKILLS

Languages: Python, JavaScript, TypeScript, Ruby, C#, D, C, SQL

Evaluation & Benchmarking: LLM Evaluation, LLM-as-a-Judge, Rubric Design, Benchmark Design, Blind Evaluation Protocols, Held-out Verifier Truth Sets, Fixture Generation, Trial Variance Analysis, Judge-vs-Human Calibration, Contamination Control

Test & Automation Infrastructure: Selenium, W3C WebDriver, Version-gated Test Selection, Coverage Aggregation, CI Matrices, RPA, Docker

Agentic Development: Devin, MCP Servers, Multi-agent Orchestration, Tool Calling, Structured Output

APIs & Reproducibility: REST, OAuth 2.0, JSON Schema, Pinned Dependencies, Version Manifests

EXPERIENCE

AI Operations Engineer

Nov 2025 – Present

Handshake

- Evaluated 4,000+ multimodal LLM responses against structured rubrics and gold-labeled QA'd samples, earning the Handshake Model Validation 2 — Expert credential.
- Eliminated 95% of manual annotation effort across 400+ runtime hours by building AutoHandshake!, a Ruby/JavaScript Selenium RPA suite with unattended execution, UI-stall recovery, and multi-tab orchestration.
- Mentored contractors on rubric interpretation, improving consistency across raters and reducing rework from misaligned quality standards.

AI Evaluation Engineer

Dec 2025 – Jun 2026

Snorkel AI

- Engineered adversarial red-team tasks for Claude and ChatGPT across 2 contracted projects, applying strict schema constraints to surface model blindspots and reasoning failure modes.
- Built Docker-based evaluation sandboxes with pinned dependency lockfiles and image tags to validate task solutions and enforce training-set integrity pre-release.

Software Engineer

Jul 2021 – Present

Independent client engagements

- Developed Warbler, an AI-assisted grading platform contracted for high-school use, with deterministic grading for objective items, schema-validated LLM rubric grading for open response, and ~95% agreement with instructor-assigned grades.
- Authored Intuit, a native LLM SDK for D with structured output, tool calling, and configurable-precision embeddings, used as Warbler's model layer.

SELECTED WORK: BENCHMARKING & EVALUATION

selenium-mcp | *Blind Benchmark Harness, D + Python*

- Built a blind benchmark harness comparing browser-automation MCP servers, with zero-foreknowledge subagents, generated fixtures, a verifier truth set held out from the agent, and a fresh server process per run.
- Measured round trips, state-recovery calls, wall clock, and bytes ingested per completed task across repeated runs, cutting round trips 39% versus playwright-mcp.
- Committed a reproducer importing the shipped capture path; it invalidated a previously cited profile table and located the content-reach boundary where inline state fails on a third-party SPA.

Selenium | *Test Infrastructure, D + Ruby + CI* | github.com/cetio/selenium

- Built suite-gated CI across 8 browser/OS workflows with version-gated test selection, verifying each commit against Chrome, Firefox, Edge, and Safari without compiling suites that do not apply.
- Wrote a Ruby coverage aggregator that executes the suite per browser, merges per-line coverage across runs, and reports per-file totals.

EDUCATION & CERTIFICATIONS

Columbus State Community College

Expected Dec 2026

AAS, Software Development

Handshake Model Validation 2 — Expert | MS Career Essentials Generative AI (NASBA)